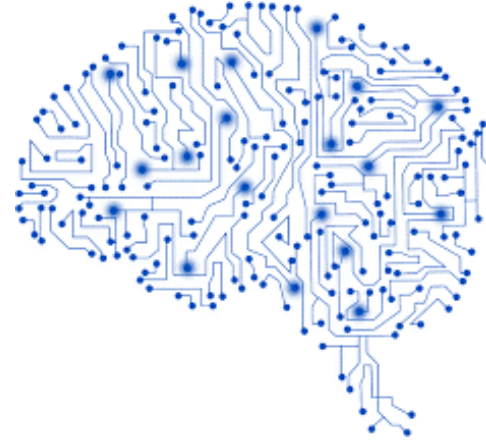




**University of Minho**  
School of Engineering



# Dados e Aprendizagem Automática

## Linear & Logistic Regression

**DAA @ MEI-1º/MiEI-4º – 1º Semestre**

Bruno Fernandes, Dalila Alves, Filipa Ferraz, Victor Alves

*Part IV*

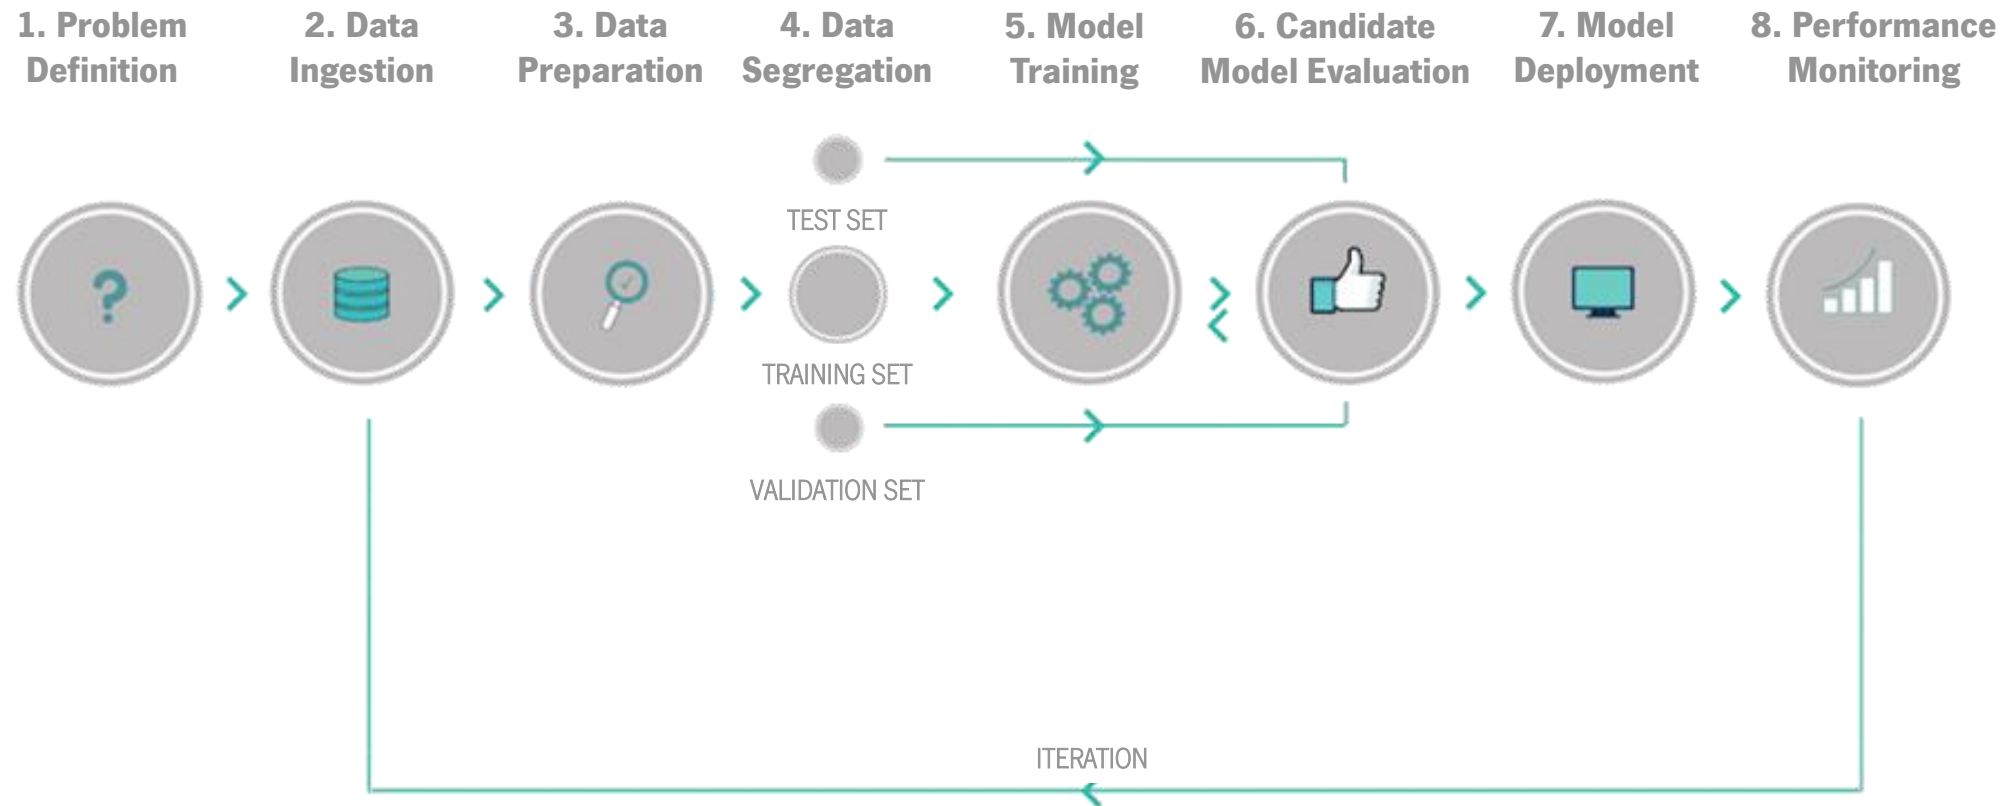
# Contents

2

- Linear Regression
- Logistic Regression
- Hands On

# Supervised Learning

3



We will explore the **Supervised Learning** algorithms.

The **choice of the algorithm** to implement depends on **the type of data and the context** provided. If we have **labelled data** or we know in advance what the **output** should be, we choose the Supervised Learning paradigm.



# Linear Regression

# The Problem and the Data

5

Problem: development of a Machine Learning model capable of **predicting house prices** for regions in the USA

Approach: **Linear Regression** approach

Dataset: table with information about houses in regions of the United States, including:

- ***Avg. Area Income***: average income of the residents of the city where the house is located
- ***Avg. Area House Age***: average age of the houses in the same city
- ***Avg. Area Number of Rooms***: average number of rooms of the houses in the same city
- ***Avg. Area Number of Bedrooms***: average number of bedrooms of the houses in the same city
- ***Area Population***: population of the city where the house is located
- ***Price***: price at which the house was sold
- ***Address***: address of the house

# Exploratory Data Analysis

6

```
USAhousing = pd.read_csv('USA_Housing.csv')
```

```
USAhousing.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386

```
USAhousing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5000 entries, 0 to 4999
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	Avg. Area Income	5000 non-null	float64
1	Avg. Area House Age	5000 non-null	float64
2	Avg. Area Number of Rooms	5000 non-null	float64
3	Avg. Area Number of Bedrooms	5000 non-null	float64
4	Area Population	5000 non-null	float64
5	Price	5000 non-null	float64
6	Address	5000 non-null	object

```
dtypes: float64(6), object(1)
```

```
memory usage: 273.6+ KB
```

# Exploratory Data Analysis

7

```
USAhousing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                      5000 non-null   float64
1   Avg. Area House Age                   5000 non-null   float64
2   Avg. Area Number of Rooms             5000 non-null   float64
3   Avg. Area Number of Bedrooms          5000 non-null   float64
4   Area Population                       5000 non-null   float64
5   Price                                 5000 non-null   float64
6   Address                               5000 non-null   object
dtypes: float64(6), object(1)
memory usage: 273.6+ KB
```

```
USAhousing.columns
```

```
Index(['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',
       'Avg. Area Number of Bedrooms', 'Area Population', 'Price', 'Address'],
      dtype='object')
```

# Exploratory Data Analysis

8

```
USAhousing.describe()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
<b>count</b>	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
<b>mean</b>	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
<b>std</b>	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
<b>min</b>	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
<b>25%</b>	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
<b>50%</b>	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
<b>75%</b>	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
<b>max</b>	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

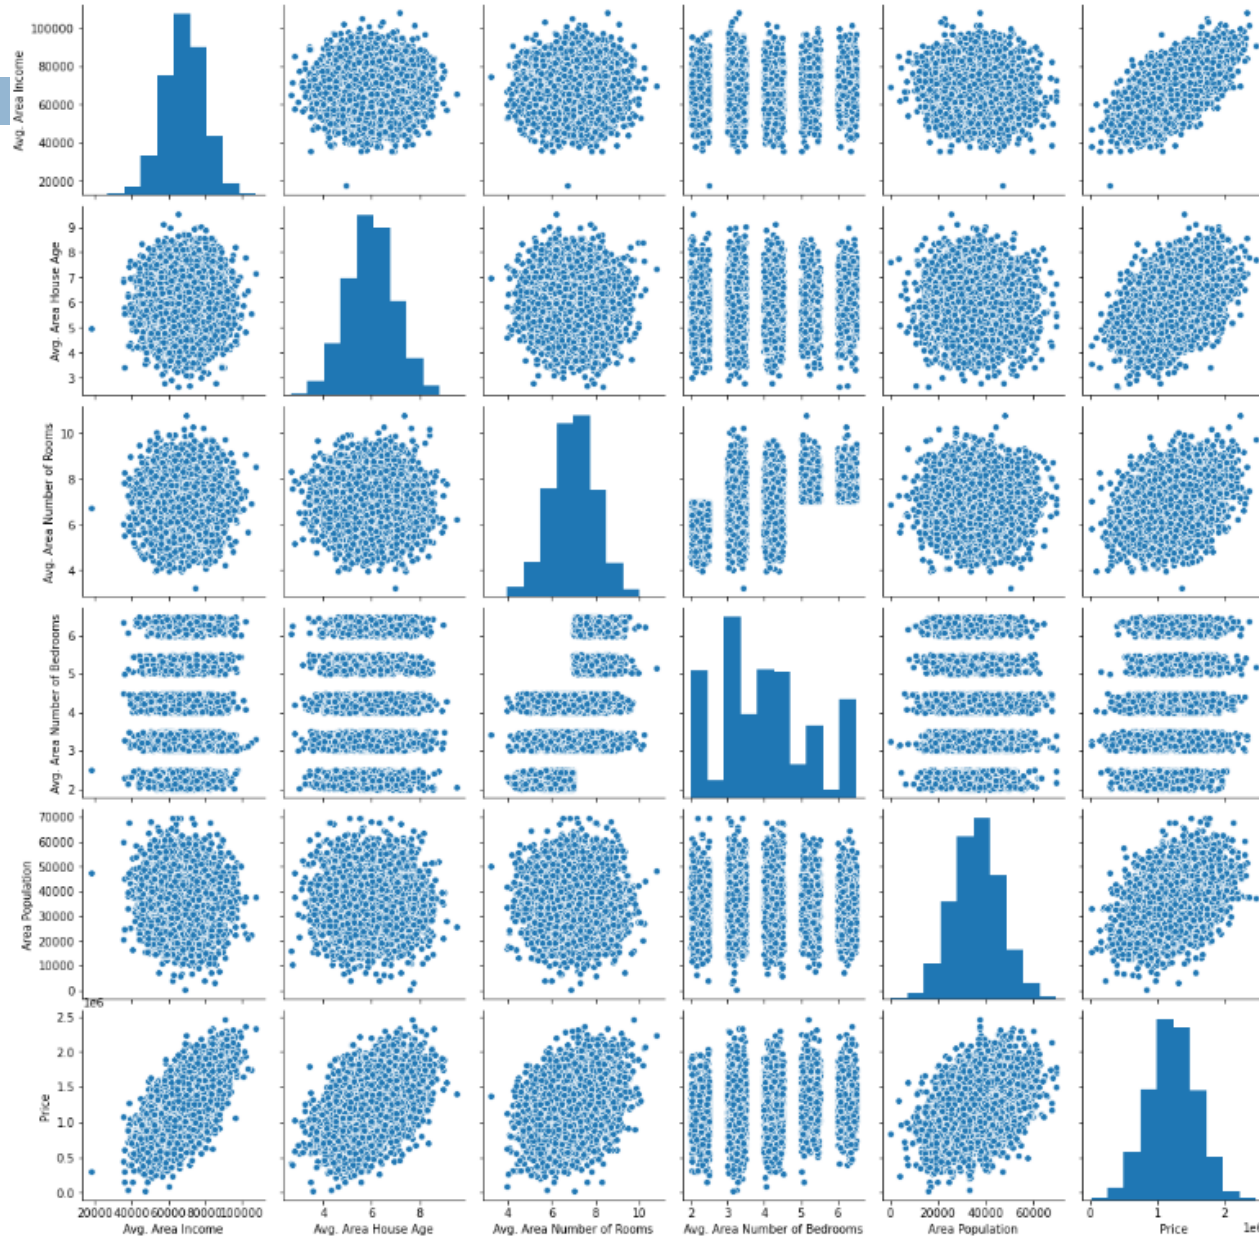


# Exploratory Data Analysis

9

Let's create some plots to check the data:


```
sns.pairplot(USAhousing, hue="Price")
```




# Exploratory Data Analysis

10

```
sns.histplot(USAhousing['Price'])
```



```
sns.heatmap(USAhousing.corr(numeric_only=True))
```



# Linear Regression

11

Let's now begin to train our model.

The target is the **Price** so we will implement a Linear Regression model.

The feature **Address** will not be considered since it is not relevant nor a numeric variable.

```
X = USAhousing[['Avg. Area Income', 'Avg. Area House Age', 'Avg. Area Number of Rooms',  
               'Avg. Area Number of Bedrooms', 'Area Population']]  
y = USAhousing['Price']
```

```
from sklearn.model_selection import train_test_split
```


```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=101)
```

# Linear Regression


12

Let's check the distribution of each subset:

```
sns.histplot(y_train)
```



```
sns.histplot(y_test)
```



# Linear Regression

13

```
sklearn.linear_model.LinearRegression(*, fit_intercept=True, copy_X=True, n_jobs=None, positive=False)
```

Creating and training the model:

```
from sklearn.linear_model import LinearRegression
```

```
lm = LinearRegression()  
lm.fit(X_train,y_train)
```

▼ LinearRegression ⓘ ?  
LinearRegression()

Then, evaluating the model by checking its coefficients and interpreting them:

```
print(lm.intercept_)
```

```
-2640159.7968525267
```

# Linear Regression

14

```
coeff_df = pd.DataFrame(lm.coef_,X.columns,columns=['Coefficient'])  
coeff_df
```

	Coefficient
Avg. Area Income	21.528276
Avg. Area House Age	164883.282027
Avg. Area Number of Rooms	122368.678027
Avg. Area Number of Bedrooms	2233.801864
Area Population	15.150420

Holding all the other features fixed:

- 1 unit increase in **Avg. Area Income** is associated with an increase of *21,53 \$*
- 1 unit increase in **Avg. Area House Age** is associated with an increase of *164883,28 \$*
- 1 unit increase in **Avg. Area Number of Rooms** is associated with an increase of *122368,68 \$*
- 1 unit increase in **Avg. Area Number of Bedrooms** is associated with an increase of *2233,80 \$*
- 1 unit increase in **Area Population** is associated with an increase of *15,15 \$*


# Linear Regression

15


Let's analyze the predictions and plot it:

```
predictions = lm.predict(X_test)
```

```
plt.scatter(y_test, predictions)
```



```
sns.histplot((y_test-predictions),bins=50);
```



# Linear Regression

16

## Regression Evaluation Metrics

The three most common evaluation metrics for regression problems are:

- Mean Absolute Error (MAE)
- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)

Comparing them:

- **MAE** is the easiest to understand because it's the *average error*,
- **MSE** is more popular than MAE because MSE *“punishes” large errors*;
- **RMSE** is even more popular than MSE because RMSE is interpretable in *units of the target variable*

All of these are **loss functions** because we want to *minimize the error*.

```
from sklearn import metrics
```

```
print('MAE:', metrics.mean_absolute_error(y_test, predictions))  
print('MSE:', metrics.mean_squared_error(y_test, predictions))  
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
MAE: 82288.22251914947
```

```
MSE: 10460958907.209057
```

```
RMSE: 102278.82922290935
```





# Logistic Regression

# The Problem and the Data

18

Problem: development of a Machine Learning model that predicts **which passengers survived** the Titanic shipwreck

Approach: **Logistic Regression** approach

Dataset: table with information regarding passengers' information, including:

- ***survival***: if the passenger survived (0: No, 1: Yes)
- ***pclass***: ticket class (1: 1<sup>st</sup>, 2: 2<sup>nd</sup>, 3: 3<sup>rd</sup>)
- ***sex***: M: Male, F: Female
- ***Age***: age in years
- ***sibsp***: number of siblings per spouses aboard
- ***parch***: number pf parents per children aboard
- ***ticket***: ticket number
- ***fare***: passenger fare
- ***cabin***: cabin number
- ***embarked***: port of embarkation (C: Cherbourg, Q: Queenstown, S: Southampton)

# Exploratory Data Analysis

19

```
train = pd.read_csv('titanic_train.csv')
```

```
train.head()
```


	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

# Exploratory Data Analysis

20

Let's check if there are missing data:

```
sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```



About 20% of the *Age* data is **missing**. The proportion of *Age* data missing is likely small enough for reasonable replacement with some form of imputation.


Looking at the *Cabin* column it looks like we are just **missing too much** of that data to do something useful with at a basic level. We will probably drop this later, or change it into another feature like “Cabin Known: 1 or 0”

# Exploratory Data Analysis


21

Let's continue on by visualizing some more of the data:

```
sns.set_style('whitegrid')  
sns.countplot(x='Survived', data=train)
```




```
sns.set_style('whitegrid')  
sns.countplot(x='Survived', hue='Sex', data=train, palette='RdBu_r')
```




# Exploratory Data Analysis

22

```
sns.set_style('whitegrid')  
sns.countplot(x='Survived', hue='Pclass', data=train, palette='rainbow')
```




```
sns.histplot(train['Age'].dropna(), kde=False, color='darkred', bins=30)
```




# Exploratory Data Analysis

23

```
train['Age'].hist(bins=30,color='darkred',alpha=0.7)
```




```
sns.countplot(x='SibSp',data=train)
```



# Exploratory Data Analysis

24

```
train['Fare'].hist(color='green',bins=40,figsize=(8,4))
```






# Data Preprocessing

25

We want to **fill in missing values** instead of dropping it. One way to do this is by filling the mean age of all passengers (**imputation**). However, we can be smarter about this and check the average age by passenger class:

```
plt.figure(figsize=(12, 7))  
sns.boxplot(x='Pclass', y='Age', data=train)
```



We can see that wealthier passengers in the higher classes tend to be older which makes sense. We will use these average age values to impute based on *pclass*.

# Data Preprocessing

26

```
def impute_age(cols):
    Age = cols[0]
    Pclass = cols[1]

    if pd.isnull(Age):

        if Pclass == 1:
            return 37


        elif Pclass == 2:
            return 29

        else:
            return 24

    else:
        return Age

train['Age'] = train[['Age', 'Pclass']].apply(impute_age,axis=1)
```

```
sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```



# Data Preprocessing

27

Now let's drop the column *cabin* and the rows in *Embarked* that are NaN:

```
train.drop('Cabin',axis=1,inplace=True)
```

```
train.head()
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	S

```
train.dropna(inplace=True)
```

# Data Preprocessing

28

We will need to convert categorical features into dummy variables using pandas' library:

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 889 entries, 0 to 890
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   PassengerId  889 non-null    int64
1   Survived     889 non-null    int64
2   Pclass       889 non-null    int64
3   Name         889 non-null    object
4   Sex          889 non-null    object
5   Age         889 non-null    float64
6   SibSp        889 non-null    int64
7   Parch        889 non-null    int64
8   Ticket       889 non-null    object
9   Fare         889 non-null    float64
10  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(4)
memory usage: 83.3+ KB
```

```
sex = pd.get_dummies(train['Sex'],drop_first=True)
embark = pd.get_dummies(train['Embarked'],drop_first=True)
```

```
train.drop(['Sex','Embarked','Name','Ticket'],axis=1,inplace=True)
```

```
train = pd.concat([train,sex,embark],axis=1)
```

```
train.head()
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	1	0	3	22.0	1	0	7.2500	True	False	True
1	2	1	1	38.0	1	0	71.2833	False	False	False
2	3	1	3	26.0	0	0	7.9250	False	False	True
3	4	1	1	35.0	1	0	53.1000	False	False	True
4	5	0	3	35.0	0	0	8.0500	True	False	True


# Logistic Regression

29


```
from sklearn.model_selection import train_test_split  
X = train.drop('Survived',axis=1)  
y = train['Survived']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=101)
```

```
sns.set_style('whitegrid')  
sns.countplot(x='Survived', data = pd.DataFrame(y_train,columns=['Survived']))
```



```
sns.set_style('whitegrid')  
sns.countplot(x='Survived', data = pd.DataFrame(y_test,columns=['Survived']))
```



# Logistic Regression

30

```
sklearn.linear_model.LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0, fit_intercept=True, intercept_scaling=1, class_weight=None, random_state=None, solver='lbfgs', max_iter=100, multi_class='auto', verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
```

## Logistic Regression' **solvers**:

- For **small datasets** *liblinear* is a good choice whereas *sag* and *saga* are faster for **larger ones**;
- For **multiclass** problems only *newton-cg*, *sag*, *saga* and *lbfgs* handle multinomial loss;
- *liblinear* is limited to **one-versus-rest schemes**.

## Supported **penalties** by solver:

- ***newton-cg*** – [*l2*, *none*]
- ***lbfgs*** – [*l2*, *none*]
- ***liblinear*** – [*l1*, *l2*]
- ***sag*** – [*l2*, *none*]
- ***saga*** – [*elasticnet*, *l1*, *l2*, *none*]

```
from sklearn.linear_model import LogisticRegression
```

# Logistic Regression

31

**Model 1:** random\_state = 2022, solver = 'newton-cg'

```
starttime = time.process_time()

logmodel1 = LogisticRegression(random_state=2022, solver='newton-cg')
print(logmodel1)
logmodel1.fit(X_train,y_train)

endtime = time.process_time()
print(f"Time spent: {endtime - starttime} seconds")

LogisticRegression(random_state=2022, solver='newton-cg')
Time spent: 0.015625 seconds

predictions1 = logmodel1.predict(X_test)
```

**Model 2:** random\_state = 2022, solver = 'lbfgs'

```
starttime = time.process_time()

logmodel2 = LogisticRegression(random_state=2022, solver='lbfgs', max_iter=800)
print(logmodel2)
logmodel2.fit(X_train,y_train)

endtime = time.process_time()
print(f"Time spent: {endtime - starttime} seconds")

LogisticRegression(max_iter=800, random_state=2022)
Time spent: 0.078125 seconds
```

**Model 3:** random\_state = 2022, solver = 'liblinear'

```
starttime = time.process_time()

logmodel3 = LogisticRegression(random_state=2022, solver='liblinear')
print(logmodel3)
logmodel3.fit(X_train,y_train)

endtime = time.process_time()
print(f"Time spent: {endtime - starttime} seconds")

LogisticRegression(random_state=2022, solver='liblinear')
Time spent: 0.0 seconds

predictions3 = logmodel3.predict(X_test)
```

# Logistic Regression

32

Let's evaluate the model using precision, recall, f1-score and confusion matrix:

```
from sklearn.metrics import classification_report, ConfusionMatrixDisplay

print("With 'newton-cg': \n", classification_report(y_test,predictions1))
print("With 'lbfgs': \n", classification_report(y_test,predictions2))
print("With 'liblinear': \n", classification_report(y_test,predictions3))
```

With 'newton-cg':

	precision	recall	f1-score	support
0	0.82	0.91	0.86	163
1	0.84	0.68	0.75	104
accuracy			0.82	267
macro avg	0.83	0.80	0.81	267
weighted avg	0.83	0.82	0.82	267

With 'lbfgs':

	precision	recall	f1-score	support
0	0.82	0.91	0.86	163
1	0.84	0.68	0.75	104
accuracy			0.82	267
macro avg	0.83	0.80	0.81	267
weighted avg	0.83	0.82	0.82	267

With 'liblinear':




	precision	recall	f1-score	support
0	0.81	0.93	0.86	163
1	0.85	0.65	0.74	104
accuracy			0.82	267
macro avg	0.83	0.79	0.80	267
weighted avg	0.82	0.82	0.81	267



# Logistic Regression

33

```
ConfusionMatrixDisplay.from_predictions(y_test, predictions1)  
ConfusionMatrixDisplay.from_predictions(y_test, predictions2)  
ConfusionMatrixDisplay.from_predictions(y_test, predictions3)  
plt.show()
```





Hands On